

Estimation of Production Functions

Aviv Nevo, University of Pennsylvania

Fall 2018

1 Introduction

Why do we care about estimating production functions (in an IO course)?

- Production functions are components in many economic models;
- Evaluate efficiency of an industry:
 - increasing returns, economies of scale
 - complementarity / substitutability between inputs
 - cost regulation, synergies and mergers
 - economics of scope
 - learning by doing
- Productivity analysis (output per index of inputs)
 - measurement of productivity differences across firms or over time
 - effects of policy (deregulation, tariffs, etc.)
 - returns of R&D (private and social)
 - returns to adoption of new technology

As a side: I have an additional reason for covering this material, especially early in the course. It is an excellent subject to motivate many very important issues in empirical micro. Unlike, many of the models we will see later, many of the models of productions functions are linear. It is important to cover many of the issues in linear models before we move on to non-linear estimation.

2 Data

The data should include:

- An output measure (typically either number of units, or value);
- Input measures (labor, capital, R&D, material, energy);
- Being able to follow the same firms over time (i.e., panel data) is very helpful for estimation.

Common sources of data are:

- Compustat: provides the annual and quarterly Income Statement, Balance Sheet, Statement of Cash Flows, and supplemental data items on most publicly held companies in North America. Financial data items are collected from a wide variety of sources including news wire services, news releases, shareholder reports, direct company contacts, and quarterly and annual documents filed with the Securities and Exchange Commission. Compustat files also contain information on aggregates, industry segments, banks, market prices, dividends, and earnings. Depending upon the data set, coverage may extend as far back as 1950 through the most recent year-end. (For more info see <http://www.kellogg.northwestern.edu/researchcomputing/compustat.htm>)
- Longitudinal Research Database (LRD): provides a plant-level database containing detailed statistics on research and development activities of US firms. The database contains detailed company-level research and development information compiled from the annual Industrial Research and Development survey starting in 1972 through 2001. (<http://www.census.gov/econ/overview/ma0800.htm>)

Other countries have similar data sets.

- Regulated (or previously regulated) industries: As part of the regulation process many firms had to report detailed cost and production data. For an example see Markiewicz, Rose and Wolfram (2004) who use data from FERC to estimate how restructuring improved efficiency in US electricity generating plants.
- Special data sets that were made available for research through a special circumstance (e.g., consulting or personal connections).

3 The Model

A typical starting point is the Cobb-Douglas model

$$Q_{it} = e^{\alpha} L_{it}^{\beta_l} K_{it}^{\beta_k} e^{u_{it}}$$

where Q_{it} is the output of plant (or firm) i at time t , L_{it} is labor (or more generally a variable input), K_{it} is capital (a quasi fixed input), u_{it} is an error term, α , β_l and β_k are parameters to be estimated.

We might include additional right-hand side variable like material, energy, different types of labor (blue/white collar), different types of capital or R&D;

The output measure might be in physical units but in most cases, since it involves aggregation across products will be measured in dollars. In some cases it will be value added rather than sales (i.e., cost of material, energy and other short-term inputs would have already been subtracted from sales).

The error term, u_{it} , includes:

- technology or management differences
- measurement errors
- variation in external factors (e.g., weather)

Taking logs we obtain

$$y_{it} = \alpha + \beta_l l_{it} + \beta_k k_{it} + u_{it}$$

where: $y_{it} \equiv \log(Q_{it})$, $l_{it} \equiv \log(L_{it})$, $k_{it} \equiv \log(K_{it})$.

4 Econometric Issues

There are several econometric issues we have to deal with in estimating the above equation.

4.1 Specification:

- Functional form? Here we will focus only on Cobb Douglas functions (following much of the applied work). There is a significant body of working examining more flexible functional forms (e.g., Fuss and McFadden)
- What does the equation mean?
 - Aggregation across products
 - Aggregation across plants/firms (or even industries)
 - Since the dependent variable is dollar amount: potentially change in pricing not productivity

4.2 Data:

- measurement error in outputs/inputs
- measurement of capital (different types, depreciation, etc.)
- measurement of labor (types, wage vs work force, etc.)

4.3 Simultaneity:

Observed inputs may be correlated with unobserved shock and therefore OLS will yield biased and inconsistent estimates.

Example 1 *Suppose we observe a cross section of firms; some are more productive (better managers); these firms might also need less labor to produce the same output (and assume they know this therefore hire less). The bottom line these firms will produce more with less labor, thus OLS will underestimate β_L .*

Example 2 *Suppose we observe the same firms over time; in a period the firm gets a higher productivity shock (which it observes) it will hire more labor. OLS will attribute all the increase in output to the change in labor, thus overestimating β_L .*

For most of what follows we will focus on the simultaneity of labor, assuming capital is pre-determined.

4.4 Selection:

Firms observed in the market are not necessarily a random draw from the population of interest. This is especially problematic in panel data, where we observe firms over time. The sample of firms that survive over time might not be random. This will introduce bias.

In what follows we will focus almost exclusively on the simultaneity issue.

5 Solutions to the Simultaneity Problem

In what follows we will focus on the variable input, labor.

Here are some of the solutions that have been offered in the literature:

5.1 Instrumental variable (in a cross section)

look for a variable that is correlated with the variable input, labor, and uncorrelated with the shock. For example, input prices.

Problems:

- input prices might not be observed;
- Might not vary by firm (in a cross section this means this cannot be used, in a panel structure basically a time dummy)
- Even if input prices vary by firm, the variation might be correlated with the error: market power in input market; matching process between firms and inputs;

5.2 Panel Data (to the rescue)

I assume you have seen most of the material in this sub-section before. I will therefore just review it quickly.

- Assume $u_{it} = \alpha_i + \varepsilon_{it}$, where α_i is a firm specific shock and ε_{it} is “white” noise. α_i is potentially correlated with the inputs (typically, we will focus on labor). A key assumption is that the part of the error that is correlated with the input is fixed over time.
- Furthermore, assume that the data include multiple observation for each firm (and are balanced).

Fixed Effects The firm specific effect is a parameter to be estimated (we place no restrictions on it besides the fact that it is fixed over time for each firm).

The model can be estimated in several ways:

Dummy variables estimate a dummy variable for each firm. This might be tough if we have many firms and a short time series for each.

First differences estimate by OLS

$$y_{it} - y_{i,t-1} = \beta_l(l_{it} - l_{i,t-1}) + \beta_k(k_{it} - k_{i,t-1}) + \varepsilon_{it} - \varepsilon_{i,t-1}$$

Note: we can also use longer differences (more on this later).

“Within” transformation

$$y_{it} - \bar{y}_i = \beta_l(l_{it} - \bar{l}_i) + \beta_k(k_{it} - \bar{k}_i) + \varepsilon_{it} - \bar{\varepsilon}_i$$

where $\bar{y}_i = \frac{1}{T} \sum_{t=1}^T y_{it}$, \bar{l}_i and \bar{k}_i are defined similarly. OLS estimation of this equation will yield identical estimates to the dummy variables approach, but will be easier to compute (and does not suffer from the incidental parameters problem).

Note: $\hat{\alpha}_i$ can be estimated by $\hat{\alpha}_i = \bar{y}_i - \beta_l \bar{l}_i - \beta_k \bar{k}_i$.

Issues:

- Is the part of the error that impacts input choice really fixed? If not the estimates might still be biased
- The “within” transformation reduces the signal to noise ratio, thus measurement error will become more of an issue and the estimates will be biased towards zero (assuming classical ME).
- Requires strict exogeneity: l_{it} must be uncorrelated with $\varepsilon_{i\tau}$ for all t and τ , since $\bar{\varepsilon}_i$ is in the error term.

5.2.1 Random Effects

An alternative approach assumes that α_i is not correlated with labor in capital and has a known (up to parameters) distribution. Therefore, in our model OLS is unbiased and results only in loss of efficiency. For unbiased estimates one can use Total (OLS on the full sample), within, between (OLS on the firm means) or GLS. GLS is the efficient estimator (if the model is correctly specified), but biased if the fixed effects model is correct. On the other hand, the within estimator is unbiased in both cases but not efficient. Thus, we can construct a Hausman test to test between the models.

Note:

- The idea behind random effects is to try to model the heterogeneity across firms. However, in a way this is not “true” heterogeneity since it is only ex-post not ex ante and it does not impact the firm behavior (in the inputs markets);
- Under the RE model there is no bias in LS estimates, in this linear static model. This will not be the case in non-linear or dynamic models.
- Relative to the FE model this model requires more assumptions and is less general.

5.3 Correlated Random Effects

This model tries to combine the FE and RE model. Like RE assume a distribution for the firm-specific shocks but allow for correlation with the inputs, like the FE model. The main problem is that we need a model for the correlation. Advantage over FE is that it requires less data, is more efficient if correct, but yields biased estimates if mis-specified. Relative to RE, deals with potential correlation.

5.3.1 Π – Matrix Approach

An approach proposed by Chamberlain (1982) that includes all the above cases, and many more as private cases. The approach has two steps. First, you recover from the data “reduced form” coefficients. Second, you impose the restrictions of the model on these coefficients.

Step 1: Suppose, for ease of exposition, that $T = 3$ and that there is only labor in the model. The idea here is to regress quantity in each period on all leads and lags.

$$\begin{bmatrix} y_{i1} \\ y_{i2} \\ y_{i3} \end{bmatrix} = \text{const} + \begin{bmatrix} \pi_{11} & \pi_{12} & \pi_{13} \\ \pi_{21} & \pi_{22} & \pi_{23} \\ \pi_{31} & \pi_{32} & \pi_{33} \end{bmatrix} * \begin{bmatrix} l_{i1} \\ l_{i2} \\ l_{i3} \end{bmatrix} + \begin{bmatrix} v_{i1} \\ v_{i2} \\ v_{i3} \end{bmatrix}$$

These are not structural equations, just a way of summarizing the data. The coefficients can be estimated using OLS.

Step 2: We now impose the model to obtain structural restrictions on the above coefficients.

Assume, for example, that $u_{it} = \alpha_i + \varepsilon_{it}$, where ε_{it} is white noise and $E(a_i | l_{i1}, l_{i2}, l_{i3}) = 0$ (as in the random effects model) then

$$\Pi = \begin{bmatrix} \pi_{11} & \pi_{12} & \pi_{13} \\ \pi_{21} & \pi_{22} & \pi_{23} \\ \pi_{31} & \pi_{32} & \pi_{33} \end{bmatrix} = \begin{bmatrix} \beta_l & 0 & 0 \\ 0 & \beta_l & 0 \\ 0 & 0 & \beta_l \end{bmatrix}$$

Alternatively assume that $E(a_i | l_{i1}, l_{i2}, l_{i3}) = \lambda_1 l_{i1} + \lambda_2 l_{i2} + \lambda_3 l_{i3}$, which is one way to interpret the fixed effects model, then

$$\Pi = \begin{bmatrix} \pi_{11} & \pi_{12} & \pi_{13} \\ \pi_{21} & \pi_{22} & \pi_{23} \\ \pi_{31} & \pi_{32} & \pi_{33} \end{bmatrix} = \begin{bmatrix} \beta_l + \lambda_1 & \lambda_2 & \lambda_3 \\ \lambda_1 & \beta_l + \lambda_2 & \lambda_3 \\ \lambda_1 & \lambda_2 & \beta_l + \lambda_3 \end{bmatrix}$$

The structural coefficients can be estimated by a Minimum Distance procedure which yields

$$\hat{\theta} = (G'V^{-1}G)^{-1} G'V^{-1} \text{vec}(\hat{\Pi})$$

where $\text{vec}(\hat{\Pi})$ is a vector containing all the coefficients estimated in Step 1, V is the covariance matrix of these estimates, $\hat{\theta}$ are the estimated structural parameters and G is the matrix of restrictions such that $\text{vec}(\Pi) = G\theta$.

For example, in the random effects model $\theta = \beta_l$, $\text{vec}(\hat{\Pi}) = [\pi_{11} \ \pi_{12} \ \pi_{13} \ \dots \ \pi_{33}]'$ and $G = [1 \ 0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \ 1]'$.

In the fixed effects model $\theta = [\beta_l \ \lambda_1 \ \lambda_2 \ \lambda_3]'$,

$$G = \begin{bmatrix} 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \end{bmatrix}.$$

Chamberlain also provides a test that can be used to test if the restrictions are “close enough” to the data.

Note: The model’s restrictions can be imposed directly and in some sense there is no need for the two stage approach. However, the 2-stage approach is appealing:

- It requires estimating the first stage only once. In some cases this might have computational advantages. It also has the advantage of only requiring a few moments from the data (which might be useful if access to the data is restricted);
- It let’s one eyeball the data and get a feel for where the model might be failing.

5.4 First Order Conditions

- This approach is based on the information in the first order conditions of optimizing firms.
- For example for a firm operating in perfectly competitive input and output markets, static cost minimization implies that

$$\frac{\partial Y}{\partial L} \frac{L}{Y} = \frac{wL}{pY} \quad \text{and} \quad \frac{\partial Y}{\partial K} \frac{K}{Y} = \frac{rL}{pY}$$

where w , r , and p are the prices of inputs and the output.

- With a Cobb-Douglas production function these output elasticities are the production function coefficients, so observations on revenue share of labor and capital can provide estimates of the coefficients

- But
 - This assume static cost minimization, i.e, no dynamics adjustment costs, etc.
 - It assumes firms are operating in a perfectly competitive output market. If this is not the case then the right hand side should be multiplied by the percent markup, $\mu = \frac{p}{mc}$. Profit maximization implies that $\mu = \frac{\eta}{1+\eta}$, where η is the elasticity of demand. If the elasticity of demand is known one could identify the production function coefficients using this approach. Alternatively, one could identify both the demand elasticity and the production function coefficients by adding additional restrictions.
 - We also assume no measurement error, or at least one of a particular form. Otherwise, simply taking shares will not give us an unbiased estimator of the coefficients.

5.5 GMM/ Differencing/Dynamic Panel

- As we saw before differencing the data (e.g., first differences) can help get rid of the additive firm-specific effect. Without measurement error and with strict exogeneity a "within" transformation is more efficient than differencing. However, if we worry about either measurement error or the strict exogeneity assumption we could consider differencing the data.

To improve efficiency, there are several possible differences one could take (for example if $T = 4$, three first- differences, two 2nd differences, one 3rd or "long" difference). Which should we take? Can we combine these?

GMM is an easy way to combine the moments. Griliches-Hausman (JOE, 1985) explore this and show that one can use the different differences to learn about the effect of measurement error.

5.5.1 Dynamic Panel Methods

This literature examines more complex models that allow for the part of the error term that is transmitted to inputs to vary over time. Now simple differencing will not solve the simultaneity problem, so the literature uses the structure of the panel to generate Instrumental Variables.

Arellano and Bond Assume that $u_{it} = \alpha_i + \omega_{it} + \varepsilon_{it}$, where
 α_i and ω_{it} are “transmitted” (i.e., impact the choice of l_{it});
 ω_{it} is not autocorrelated;

Taking differences we get

$$y_{it} - y_{i,t-1} = \beta_l(l_{it} - l_{i,t-1}) + \beta_k(k_{it} - k_{i,t-1}) + (\omega_{it} - \omega_{i,t-1}) + (\varepsilon_{it} - \varepsilon_{i,t-1})$$

The difference $(\omega_{it} - \omega_{i,t-1})$ is correlated with $(l_{it} - l_{i,t-1})$, therefore we need an IV. Arellano and Bond (ReStud, 1991) suggest using lagged values of output and the inputs as IV (i.e., use $l_{i,t-\tau}$ and $k_{i,t-\tau}$, for $\tau \geq 2$). We could also consider using lagged values of the outputs (i.e., $y_{i,t-\tau}$ for $\tau \geq 2$) and depending on the assumptions on how/when capital is chosen $k_{i,t}$ and $k_{i,t-1}$, as additional IVs. The estimation is based on the conditional moment

$$E \left[u_{it} - u_{i,t-1} \mid (l_{i\tau}, k_{i\tau})_{\tau=1}^{t-2} \right] = 0 \quad (1)$$

Notes:

- In practice this approach has performed poorly both with real data and in Monte Carlo studies.
- The reason seems to be that lagged values are weak instruments for the differences.
- The assumption that ω_{it} is not autocorrelated seems strong;

Blundell and Bond (JOE, 1998, Econometric Reviews, 1999) propose getting around these problems.

To deal with the weak IV problem they propose additional moment conditions. In particular, they propose using lagged *differences* of the inputs (i.e., $l_{it-1} - l_{i,t-2}$, and $k_{it-1} - k_{i,t-2}$) as instruments for the *levels* equation. This is justified if we assume

$$E \left[u_{it} \mid \{l_{i\tau} - l_{i\tau-1}, k_{i\tau} - k_{i\tau-1}\}_{\tau=2}^{t-2} \right] = 0. \quad (2)$$

In words we assume that the level of past inputs is not a valid IV (since it will depend on α_i), but that past changes in inputs are valid (since they were driven

by realizations of past ω and ε). These IVs are likely to be powerful because there are some adjustment costs and the current level of the inputs will likely depend on past levels, which are a summation of past changes. A natural question is if there are adjustment costs then why aren't firms forward looking, in which case if they can (partially) predict future ω and ε the moment condition in (2) will be violated. We'll return to this below.

Combining the "levels as IVs for difference equations" moments in (1), with the "differences as IVs for levels equations" in (2) yields what is often called the "system GMM". This typically has better, more stable, performance than just relying on either the moments in (1) or (2) separately.

These ideas can be extended to deal with the case of serial correlation in ω_{it} .

Assume that $u_{it} = \alpha_i + \omega_{it} + \varepsilon_{it}$, where

α_i and ω_{it} are "transmitted" (i.e., impact the choice of l_{it});

$\omega_{it} = \rho\omega_{it-1} + v_{it}$ AR(1) process;

B-B propose using $y_{i\tau}$, $l_{i,t-\tau}$ and $k_{i,t-\tau}$, for $\tau \geq 3$ as IV for the (quasi-difference) equation

$$(y_{it} - \rho y_{i,t-1}) - (y_{it-1} - \rho y_{i,t-2}) = \beta_l ((l_{it} - \rho l_{i,t-1}) - (l_{it-1} - \rho l_{i,t-2})) \\ + \beta_k ((k_{it} - \rho k_{i,t-1}) - (k_{it-1} - \rho k_{i,t-2})) + \varepsilon_{it}^*$$

where

$$\varepsilon_{it}^* = v_{it} - v_{it-1} + (\varepsilon_{it} - \rho \varepsilon_{i,t-1}) - (\varepsilon_{it-1} - \rho \varepsilon_{i,t-2})$$

The estimation is based on the conditional moment

$$E \left[\varepsilon_{it}^* \mid (y_{i\tau}, l_{i\tau}, k_{i\tau})_{\tau=1}^{t-3} \right] = 0$$

Notes:

- The idea here can be extended to higher order *linear* processes, but generally not to non-linear process (first, or higher order);

- The exact lags used for the estimation can be modified slightly depending on the assumption we want to make.
- The estimation can be performed in a several ways. We can generate moment conditions, which are non-linear in the parameters, by interacting the error term, ε_{it}^* , with the IVs. The parameters can then be estimated using GMM. This is the same as estimating the equation

$$y_{it} - y_{it-1} = \rho(y_{i,t-1} - y_{i,t-2}) + \beta_l(l_{it} - l_{it-1}) + \rho\beta_l(l_{i,t-1} - \rho l_{i,t-2}) \\ + \beta_k(k_{it} - k_{it-1}) + \rho\beta_k(k_{i,t-1} - k_{i,t-2}) + \varepsilon_{it}^*$$

The parameters can also be estimated in a two step MD procedure. First, we estimate the above equations without imposing the restrictions of the model (i.e., running 2SLS of $y_{it} - y_{it-1}$ on $y_{i,t-1} - y_{i,t-2}$, and 2 lags of the differences in l and k) and then imposing the restrictions in a second stage Minimum Distance.

- As before, we can also add the differences as IVs for levels equations as in equation (2).

5.6 A "Structural" Approach: Olley Pakes (The Dynamics of Productivity in the Telecommunications Equipment Industry)

One of the main problems with the dynamic panel data approach discussed above is that the setup is mostly statistical in nature, with little connection to economic modeling. (Indeed, I spared you most of the discussion of the exact statistical assumptions required in all these cases). An alternative approach that is closer tied to economic theory and therefore often considered more "structural" is offered by Olley-Pakes. As we will see, this approach has a lot in common with the dynamic panel methods.

The paper studies the effect of deregulation on productivity. It does so by estimating a production function and using it to recover productivity. The paper deals with simultaneity and selection, but I will focus only on the primer. The data use is LRD data, which we discussed earlier.

The basic idea uses ideas somewhat similar to the FOC approach, namely, to bring in new information on the choice of inputs. The difference is that there is not going to be a reliance on an exact model of static profit maximization in perfectly competitive markets. Instead the relationship will be imposed non-parametrically and inverted. The inversion will give us an estimate of the unobserved productivity that we will use as a control.

5.6.1 The model

$$y_{it} = \alpha + \beta_l l_{it} + \beta_k k_{it} + \omega_{it} + \varepsilon_{it}$$

where ω_{it} is transmitted into the labor choice and ε_{it} is white noise. Note: 1) The paper also includes the age of the plant, a_{it} ; 2) Unlike before, ω_{it} is not fixed over time.

Assumption 1 (first-order Markov): $P(\omega_{it} \mid I_{it-1}) = P(\omega_{it} \mid \omega_{it-1})$, where I_{it-1} is the firm's information set at time $t - 1$.

This assumption is more general than the linear AR(1) typically assumed in the dynamic panel literature. However, it rules out higher order persistence, for example as in Arellano and Bond, where the shock includes a firm "fixed" effect.

Assumption 2 (timing): (i) labor is a non-dynamic input (choice of labor at time t does not impact future profits); (ii) capital choice is dynamic and evolves according to $k_{it} = K(k_{it-1}, i_{it-1})$, where $K()$ is a deterministic function and i_{it-1} is investment at time $t - 1$.

The second part of the assumption (together with Assumption 1) implies that $\omega_{it} - E(\omega_{it} \mid I_{it-1}) = \omega_{it} - E(\omega_{it} \mid \omega_{it-1})$ is not correlated with k_{it} (since i_{it-1} is determined at $t - 1$).

Assumption 3 (strict monotonicity): $i_{it} = f_t(\omega_{it}, k_{it})$ is strictly increasing in ω_{it} .

Note

- labor is non-dynamic so it does not enter the investment function
- only ω_{it} , no other unobserved shocks, enter the function. An example of such shocks are firm specific input prices (or a firm "fixed" effect).
- f is indexed by t to allow for changes in market conditions (e.g., the macro economy) that change over time and impact all firms.

Assumption 3 allows us to invert the investment decision to recover the unobservable:

$$\omega_{it} = f_t^{-1}(i_{it}, k_{it}).$$

So if we know f_t^{-1} , we could control for ω_{it}

$$y_{it} = \alpha + \beta_l l_{it} + \beta_k k_{it} + f_t^{-1}(i_{it}, k_{it}) + \varepsilon_{it} \equiv \beta_l l_{it} + \phi_t(i_{it}, k_{it}) + \varepsilon_{it}$$

Since we do not know f_t^{-1} (we need to solve a fairly complex dynamic problem in order to get it), we will estimate it non-parametrically.

5.6.2 Estimation

Step 1 Estimate β_l by regression of y_{it} on l_{it} and a non-parametric estimate of ϕ_t , get $\hat{\beta}_l$ and $\hat{\phi}_{it}$.

(For the purpose of the discussion here you can think of a non-parametric estimate as a polynomial expansion of pre-set order, say 4th order).

Step 2 (recovering β_k) We note that

$$\omega_{it} = E(\omega_{it} \mid I_{it-1}) + \xi_{it} = E(\omega_{it} \mid \omega_{it-1}) + \xi_{it}$$

where ξ_{it} is the unexpected innovation in ω_{it} . The second equality follows from Assumption 1.

By construction $E(\xi_{it} \mid I_{it-1}) = 0$, which combined with Assumption 2 implies that

$$E(\xi_{it} \mid k_{it}) = 0$$

How can we use this?

The Step 1 estimates allow us to write $\omega_{it}(\beta_k) = \hat{\phi}_{it} - \beta_k k_{it}$ and therefore

$$\xi_{it}(\beta_k) = \omega_{it}(\beta_k) - \psi(\omega_{it-1}(\beta_k))$$

where $\psi(\omega_{it-1}) \equiv E(\omega_{it} \mid \omega_{it-1})$. Since we allow for a general (unknown) first-order Markov process ψ is unknown and needs to be estimated, which we will do non-parametrically. If we assumed an AR(1) process ψ will just be a linear function.

This can be used in several ways. For example, we could choose $\hat{\beta}_k$ (and the parameters of the polynomial expansion of $\psi(\cdot)$) to

$$\text{Min} \sum_{it} \xi_{it}^2(\beta_k)$$

Alternatively, we can construct a moment estimator by:

1. for a given β_k ; computing $\omega_{it}(\beta_k)$;
2. non-parametrically regressing ω_{it} on ω_{it-1} to get estimates of ψ ;
3. use these estimates to compute ξ_{it} and (iv) search over β_k to satisfy the sample analog of the above moment condition $\frac{1}{T} \frac{1}{N} \sum_{it} \xi_{it}(\beta_k) k_{it} = 0$.

Of course, we could stack up all the moments implied by this procedure and solve this in one step.

In summary, we identified

- β_l by controlling for ω_{it} using the investment decision, which brought in new information. Since, investment is closely related to the change in capital in some ways this parallels the ideas in using lagged differences to estimate a levels equation employed in the dynamic panel literature. The difference is that here the lagged difference are used as a "control function" of sort.
- β_k is identified off of the timing decisions, which very much parallels the ideas in using past levels to identify a (quasi) difference equation in the dynamic panel literature.

While we talked about the estimation as been done in two steps, in reality this can be done in one step, stacking up all the moments.

Note, the model in the paper is more general in two ways:

- They include age in the model. Does not change much.
- Allow for selection, which depends on ω_{it} . This adds a step, between Step 1 and 2, where the probability of exit is estimated non-parametrically. Denote this probability by \widehat{p}_{it} . Step 2 is then slightly modified by writing $\psi(\omega_{it-1}, \widehat{p}_{it})$

5.6.3 Levinsohn-Petrin

In many data sets investment is lumpy. In particular many observations with $i_{it} = 0$.

For these observations f_t is not strictly increasing and cannot be inverted.

In principle, the OP approach can still work but it requires that we ignore all the observations with $i_{it} = 0$, which could mean ignoring a lot of data. For this reason Levinsohn-Petrin propose to modify the OP approach.

The consider

$$y_{it} = \alpha + \beta_l l_{it} + \beta_k k_{it} + \beta_m m_{it} + \omega_{it} + \varepsilon_{it}$$

where m_{it} are intermediate inputs like material, fuel or electricity.

Assume that $m_{it} = f_t(\omega_{it}, k_{it})$ and the f_t is invertible.

Therefore

$$\omega_{it} = f_t^{-1}(m_{it}, k_{it}).$$

The estimation follows OP very closely.

In Step 1, regress y_{it} on l_{it} and a non-parametric estimate of $\phi_t(m_{it}, k_{it})$, get $\hat{\beta}_l$ and $\hat{\phi}_{it}$.

In Step 2 exploit the conditional moment condition:

$$E(\xi_{it}(\beta_k, \beta_m) \mid k_{it}, m_{it-1}) = 0$$

Comments:

- Nice idea to bring in more proxies (allows for test, and possibly a richer model);
- Why is f_t nonparametric? Unlike OP we do not need a dynamic model to specify it. Given the production function and assumptions about input and output prices it could be specified. See a recent paper by Gandhi et al.
- In some sense no new information - what is unique about m_{it} ? Why couldn't we use labor?
- Timing assumptions.

5.6.4 Akerberg-Caves-Frazer

They critique the LP (and to a lesser extent the OP) method. Furthermore, they offer an alternative. I'll just quickly summarize their main criticism and review their alternative.

"Colinearity" Issues Consider Step 1 in LP

$$y_{it} = \beta_l l_{it} + \phi_t(m_{it}, k_{it}) + \varepsilon_{it}$$

The question is why (in the model) will l_{it} vary independently of $\phi_t(m_{it}, k_{it})$? Obviously, they vary independently when we estimate the model (otherwise, we could not get first stage estimates). But the question is this real variation or just a function of us not allowing for a general enough function for ϕ_t . (or of the model being mis-specified)

ACF consider various options on timing.

1) l_{it} chosen at the same time as m_{it} in this case, just like the way we described m_{it} we can write

$$l_{it} = g_t(\omega_{it}, k_{it}) = g_t(f_t^{-1}(m_{it}, k_{it}), k_{it}) = h_t(m_{it}, k_{it})$$

No independent variation in l_{it}

2) l_{it} chosen either before or after m_{it} (and ω_{it} evolves between the choices)
If l_{it} is chosen after m_{it} , we will get variation in l_{it} (because it includes additional information), but inverting m_{it} will not recover the correct productivity shock.

If l_{it} is chosen before m_{it} , then $m_{it} = f_t(l_{it}, \omega_{it}, k_{it})$ and we will not get the required variation

3) A couple of assumptions that might work measurement error in l_{it} , but not in m_{it} will generate the required variation and still allow the inversion to recover the productivity shock (note that we need no measurement error in m_{it} for the inversion to work).

Another alternative is that l_{it} is chosen after m_{it} , ω_{it} does *not* evolve between the choices, and there is an additional unexpected error that impacts the choice of labor.

The last two alternatives will work but are very specific and it is not clear we want to build on them to justify the procedure

Similar argument can be made for OP (see the ACF paper for details). The bottom line is that similar issues exist in OP, but that the assumptions required to justify the procedure are easier to believe.

Alternative method In addition to pointing out the problem, ACF offer an alternative method (they actually offer more than one, but I will discuss the main one).

Consider the value added production function:

$$y_{it} = \alpha + \beta_l l_{it} + \beta_k k_{it} + \omega_{it} + \varepsilon_{it}$$

The timing is as follows:

- i_{it} (and k_{it}) set at time $t - 1$;
- l_{it} set at time $t - b$, $0 < b < 1$;
- m_{it} set at time t ; and the production occurs.

The productivity shocks evolves according to $P(\omega_{it} \mid I_{it-b}) = P(\omega_{it} \mid \omega_{it-b})$ and $P(\omega_{it-b} \mid I_{it-1}) = P(\omega_{it-b} \mid \omega_{it-1})$.

Finally, given the timing materials are determined by $m_{it} = f_t(l_{it}, \omega_{it}, k_{it})$, and f_t is invertible.

Therefore:

$$y_{it} = \alpha + \beta_l l_{it} + \beta_k k_{it} + f_t^{-1}(l_{it}, m_{it}, k_{it}) + \varepsilon_{it}$$

The estimation is in two steps

Step 1, regress y_{it} on a non-parametric function of l_{it}, k_{it} and m_{it} . Note, that unlike before no parameters are estimated only $\widehat{\phi}_{it}$ is recovered (basically the whole purpose of this step is to net out ε_{it}).

Step 2, estimate the coefficients β_l and β_k using the conditional moment

$$E(\xi_{it}(\beta_l, \beta_k) \mid k_{it}, l_{it-1}) = 0$$

where as before $\xi_{it} = \omega_{it} - E(\omega_{it} \mid \omega_{it-1})$.

The advantage of this approach is that it makes assumptions about the timing that seem reasonable and are consistent with the estimation. Note, that we did not really need materials: we could have used labor as the proxy.

5.6.5 Final Comments

- The OP/LP/ACF estimators can also be estimated in a single step, by stacking up all the relevant moments. the advantage of a one step approach is in computing the standard errors.
- Relation to the dynamic panel literature
 - OP/LP/ACF allow for general first order process (not just a linear AR(1))
 - dynamic panel approach can allow for a fixed effect in addition to the AR(1) process (i.e., more persistence in the productivity shock), while OP/LP/ACF cannot. If a fixed-effect exists the moment condition used in the second step would not be valid.
 - dynamic panel approach does not require scalar unobservable or monotonicity condition.

6 Results

Olley-Pakes Griliches-Mairesse Akerberg-Caves-Frazer

TABLE VI
ALTERNATIVE ESTIMATES OF PRODUCTION FUNCTION PARAMETERS^a
(STANDARD ERRORS IN PARENTHESES)

Sample:	Balanced Panel		Full Sample ^{c,d}						
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	Nonparametric F_w	
Estimation Procedure	Total	Within	Total	Within	OLS	Only P	Only h	Series	Kernel
Labor	.851 (.039)	.728 (.049)	.693 (.019)	.629 (.026)	.628 (.020)			.608 (.027)	
Capital	.173 (.034)	.067 (.049)	.304 (.018)	.150 (.026)	.219 (.018)	.355 (.02)	.339 (.03)	.342 (.035)	.355 (.058)
Age	.002 (.003)	-.006 (.016)	-.0046 (.0026)	-.008 (.017)	-.001 (.002)	-.003 (.002)	.000 (.004)	-.001 (.004)	.010 (.013)
Time	.024 (.006)	.042 (.017)	.016 (.004)	.026 (.017)	.012 (.004)	.034 (.005)	.011 (.01)	.044 (.019)	.020 (.046)
Investment	—	—	—	—	.13 (.01)	—	—	—	—
Other Variables	—	—	—	—	—	Powers of P	Powers of h	Full Polynomial in P and h	Kernel in P and h
# Obs. ^b	896	896	2592	2592	2592	1758	1758	1758	1758

^aThe dependent variable in columns (1) to (5) is the log of value added, while in columns (6) to (10), the dependent variable is the log of value added $-b_l \cdot \log(\text{labor})$.

^bThe number of observations in the balanced panels of regressions 1 and 2 are the observations for those plants that have continuous data over the period, with zero investment observations removed. The 2592 observations used in columns (3), (4), and (5) are all observations in the full sample except those with zero investment. Approximately 8% of the full data set had observations with zero investment. Columns (6) to (10) have fewer observations because the sampling procedures for the Annual Survey of Manufactures forced us to drop observations in years 1978, 1983, and the last year, 1987. See note c.

^cThe number of observations in the last four columns decreases to 1758 because we needed lagged values of some of the independent variables in estimation. This rules out using the first observation on each plant and the first year of the rotating five-year panels that make up the Annual Survey of Manufactures. To check that the difference between the estimates in columns (6)–(9) and those in columns (3)–(5) are not due to the sample, we ran the estimating equations in columns (3)–(5) on the 1758 plant sample and got almost identical results.

^dConsult the text for details of the estimation algorithm for columns (6) to (10).

Table 3: Alternative Estimates of Production Function Parameters¹:
U.S. R&D Performing Firms, 1973, 1978, 1983, 1988
(standard errors in parentheses)

Variables ¹	Sample					
	Balanced Panel		Full Sample ³			
	(1)	(2)	(3)	(4)	(5)	(6)
	Total	Within	Total OLS		Nonparametric F	
Labor	.496 (.022)	.685 (.030)	.578 (.013)	.551 (.013)	.591 (.013)	
Physical Capital	.460 (.014)	.180 (.027)	.372 (.009)	.298 (.012)	.321 (.016)	.320 (.017)
R&D Capital	.034 (.015)	.099 (.027)	.038 (.007)	.027 (.007)	.081 (.016)	.077 (.019)
Investment	-	-	-	.110 (.011)	-	
Other Variables ⁴	-	-	-	-	Powers of h	Polynomial in P and h
# Observations ²	856		2971		1571	

(1) The dependent variable in columns 1 to 4 is the log of sales, while in column 5 and 6, the dependent variable is the $\log(\text{value added}) - \beta \cdot \log(\text{labor})$.

(2) The number of observations in the balanced panel for regressions in columns 1 and 2 are the observations for those firms that have continuous data over the period. Similarly, the 2971 observations in columns 3 and 4 are all observations in the full sample. (Only six observations had to be discarded because of zero investment.) The number of observations in the last two columns (5) and (6) decreased to 1571 because lagged values of some of the independent variables are needed in estimation.

(3) Consult the text for details of the estimation algorithm leading up to columns 5 and 6.

(4) The other variables in the equations are: Year, and Year x Industry 357 (i.e. computers) dummy variables.

TABLE I
MONTE CARLO RESULTS^a

Meas. Error	ACF				LP			
	β_l		β_k		β_l		β_k	
	Coef.	Std. Dev.	Coef.	Std. Dev.	Coef.	Std. Dev.	Coef.	Std. Dev.
<i>DGP1—Serially Correlated Wages and Labor Set at Time $t - b$</i>								
0.0	0.600	0.009	0.399	0.015	0.000	0.005	1.121	0.028
0.1	0.596	0.009	0.428	0.015	0.417	0.009	0.668	0.019
0.2	0.602	0.010	0.427	0.015	0.579	0.008	0.488	0.015
0.5	0.629	0.010	0.405	0.015	0.754	0.007	0.291	0.012
<i>DGP2—Optimization Error in Labor</i>								
0.0	0.600	0.009	0.400	0.016	0.600	0.003	0.399	0.013
0.1	0.604	0.010	0.408	0.016	0.677	0.003	0.332	0.011
0.2	0.608	0.011	0.410	0.015	0.725	0.003	0.289	0.010
0.5	0.620	0.013	0.405	0.017	0.797	0.003	0.220	0.010
<i>DGP3—Optimization Error in Labor and Serially Correlated Wages and Labor Set at Time $t - b$ (DGP1 plus DGP2)</i>								
0.0	0.596	0.006	0.406	0.014	0.473	0.003	0.588	0.016
0.1	0.598	0.006	0.422	0.013	0.543	0.004	0.522	0.014
0.2	0.601	0.006	0.428	0.012	0.592	0.004	0.473	0.012
0.5	0.609	0.007	0.431	0.013	0.677	0.005	0.386	0.012

^a1000 replications. True values of β_l and β_k are 0.6 and 0.4, respectively. Standard deviations reported are of parameter estimates across the 1000 replications.

TABLE II
MONTE CARLO RESULTS^a

Meas. Error	OP			
	β_l		β_k	
	Coef.	Std. Dev.	Coef.	Std. Dev.
<i>DGP1—Serially Correlated Wages and Labor Set at Time $t - b$</i>				
0.0	0.824	0.002	0.188	0.004
0.1	0.836	0.002	0.179	0.004
0.2	0.845	0.002	0.171	0.003
0.5	0.865	0.002	0.151	0.003